

# Statistics 360: Advanced R for Data Science

## Lecture 6

Brad McNeney

Debugging

Measuring performance

# Debugging (Ch22) and Measuring performance (Ch 23)

- ▶ Reading: Text, Chapters 22 and 23
- ▶ Topics on debugging:
  - ▶ overview of debugging
  - ▶ tracing execution with `traceback()`
  - ▶ interactive debugging with `debug()` and `browser()`
  - ▶ non-interactive debugging: `dump.frames()` and printing
  - ▶ test cases to detect future bugs
- ▶ Topics on measuring performance:
  - ▶ profiling
  - ▶ microbenchmarking
  - ▶ final thoughts

# Debugging

# Overview

- ▶ Focus on the easy part of debugging: finding and fixing the source of unexpected errors.
  - ▶ Mistakes that give incorrect results but throw no errors are harder to find.
- ▶ Workflow tips for finding and fixing errors
  - ▶ Google it: If you don't understand the error message, try pasting it into a google search.
  - ▶ Make a small self-contained example (reproducible example, a.k.a. `reprex`).
  - ▶ Find it with tools like `traceback()`, `debug()` and `browser()`.
  - ▶ Fix it and make a test case to alert you if you accidentally re-introduce the bug.

# Reproducible examples

- ▶ Reproducible means including any source code, data and library calls so that the code can run as it did when the error was triggered.
- ▶ Next reduce the code to a minimal example that triggers the problem.
  - ▶ For example, remove lines of code, compute on a smaller R object, use build-in data.
- ▶ The act of creating the reprex may show you the error.
- ▶ If not, you are in a position to ask for help from a class-mate, mailing list or stack overflow.
- ▶ I find it hard to construct reprexs without first finding the lines that throw the error ...

# Tracing execution

- ▶ After an error, you can use `traceback()` to see the sequence of function calls (“call stack”) that lead to the error.
  - ▶ The numbers in each entry are supposed to be line numbers of the call in the calling function, but they usually just confuse me

...

```
f <- function(x) { g(h(x)) }  
g <- function(x) {  
  x  
}  
h <- function(x) {  
  if(!is.numeric(x)) stop("x must be numeric")  
}  
# f("cat") # uncomment to run  
# traceback()
```

# Interactive debugging

- ▶ Main tools are `browser()` and `debug()`.
- ▶ Stop and step through function execution.
  - ▶ Can print variables and execute R commands to investigate

```
h <- function(x) {  
  browser()  
  if(!is.numeric(x)) stop("x must be numeric")  
}  
#f("cat")
```



## browser commands

- ▶ `n` executes the next step. Use `print(n)` to print a variable named `n`.
- ▶ `s` is like `n` but will step into a function call.
- ▶ `f` finishes execution of the current loop or function.
- ▶ `c` leaves interactive debugging and continues regular execution.
- ▶ Enter (Return) repeats the last browser command
- ▶ `Q` completely exits the function.

## debug()

```
h <- function(x) {  
  if(!is.numeric(x)) stop("x must be numeric")  
}  
  
# debug(f)  
# f("cat")  
# undebug(f)  
# debug(g)  
# f("cat")  
# undebug(g)  
# debug(h)  
# f("cat")  
# undebug(h)
```

# Non-interactive debugging

- ▶ You can insert `print()` or `cat()` statements to see values of variables in your code if you find the trace too confusing and browser too time-consuming.

## Test cases

- ▶ After you find and fix a bug it is a good idea to devise a test of your code that will flag the problem if you ever accidentally re-introduce it.
- ▶ If you are writing an R package you should investigate the `testthat` package, which helps you compile and run “unit” tests on small pieces of your code.

```
f <- function(x) { x + 3 }  
# test  
f(3) # should return 6
```

```
## [1] 6
```

## Measuring performance

# Measuring performance

- ▶ When you write code you develop an intuition about what parts will run slowly – don't trust this!
- ▶ As statisticians we know that the only thing you can trust is data.
- ▶ Profiling and benchmarking are ways to collect data on your code

```
library(profvis) #visualize profiling data  
library(bench)  # benchmarking tools
```

# Profiling

- ▶ R uses a statistical profiler that records the call stack at small intervals.
  - ▶ Read the call stack from right to left

```
f <- function() {pause(0.1);g();h()} # pause() is from profvis
g <- function() {pause(0.1);h()}
h <- function() {pause(0.1)}
Rprof()
f()
```

```
## NULL
```

```
Rprof(NULL) # Now view Rprof.out
```

```
sample.interval=20000
```

```
"pause" "f"
```

```
"pause" "f"
```

```
"pause" "f"
```

```
"pause" "f"
```

```
"pause" "f"
```

```
"pause" "g" "f"
```

```
"pause" "g" "f"
```

```
"pause" "g" "f"
```

```
"pause" "g" "f"
```

```
"pause" "g" "f"
```

```
"pause" "h" "g" "f"
```

```
"pause" "h" "g" "f"
```

```
"pause" "h" "g" "f"
```

```
"pause" "h" "g" "f"
```

```
"pause" "h" "g" "f"
```

```
"pause" "h" "f"
```

```
"pause" "h" "f"
```

```
"pause" "h" "f"
```

```
"pause" "h" "f"
```

```
"pause" "h" "f"
```



# Summary of profile

```
# summaryRprof() # uncomment and run
```

## Visualize profile

```
source("lec6profiling.R") # profiler will refer to this source file  
# profvis({ f() }) # Or choose Profile from RStudio
```

# Microbenchmarking

# Final thoughts

- ▶ Avoid the temptation to let performance considerations dominate your code development and lead you to profile and benchmark extensively.
- ▶ Remember that the most important performance improvement is code that gives the right answer.